HOUSE PRICE PREDICTION

Machine Learning Project - Phase 2

Submitted by:

Vishvesh Rao (AM.EN.U4CSE19161) S5 CSE B

ML PROBLEM DESCRIPTION

Formal description

- Task (T): Predict the price of house
- Experience(E): A dataset which has the prices of houses from various in a particular area
- Performance(P): The number of house prices predicted correctly out of all the prices given

Assumptions

- This assumes that external factors such as weather, crime rates etc do not affect housing prices
- This model also does not take into account the location of the houses such as distance to the nearest form of public transportation or any other type of establishments.
- The data provided to us is from 1990 ranging to 2019 so we would be making prediction using the trends found in that dataset

DATASET

This Dataset is the Boston Housing Prices, Boston is a suburb in USA.

The data was drawn from the Boston Standard Metropolitan Statistical Area (SMSA) in 1990. The data pertains to the houses found in a given California district and some summary stats about them based on the 1990 census data. It consistis of 20640 instances

The attributes are defined as follows:

- 1. **longitude**: A measure of how far west a house is; a higher value is farther west
 - 2. **latitude**: A measure of how far north a house is; a higher value is farther north
 - 3. **housingMedianAge**: Median age of a house within a block; a lower number is a newer building
 - 4. **totalRooms**: Total number of rooms within a block
 - 5. **totalBedrooms**: Total number of bedrooms within a block
 - 6. **population**: Total number of people residing within a block
 - 7. **households**: Total number of households, a group of people residing within a home unit, for a block
 - 8. **medianIncome**: Median income for households within a block of houses (measured in tens of thousands of US Dollars)
 - 9. **medianHouseValue**: Median house value for households within a block (measured in US Dollars)
 - 10. oceanProximity: Location of the house w.r.t ocean/sea

PRE-PROCESSING

Initial Dataset had 8 features (excluding target feature) and consisted of 20,640 entries/samples.

Pre-processing steps

1. Feature analysis

checked all the features for null value and number of distinct values present in a feature based on these criterias removed a feature till there were no null values and all features contained a good amount of distinct values

Also checked the correlations between the features and plotted a corelation matrix.

2. Feature Modifications

Certain features were not useful/useable in their present state. The "Longitude" and "Latitude" attributes were not easily understandable so those were used to get exact "Road" and "county" names of each data sample and these and these two attributes were added to the dataframe after removing "Longitude" and "Latitude".

3. *Filing Missing Values*

The "road" and "county" attributes had numerous missing values after geoplotting of the co-ordinates was done. Consequently Logistic Regression specifically SGDC (Stochastic Gradient Descent Classifier) was used to classify the missing road and county values.

All above processes are explained in detail with pictorial/graphical representaion of data in Phase-1 Document

DATA SUMMARIZATION AND VISUALIZATION

Here we will look at each of the individual attributes in detail analysing various statistical parameters and bar graphs.

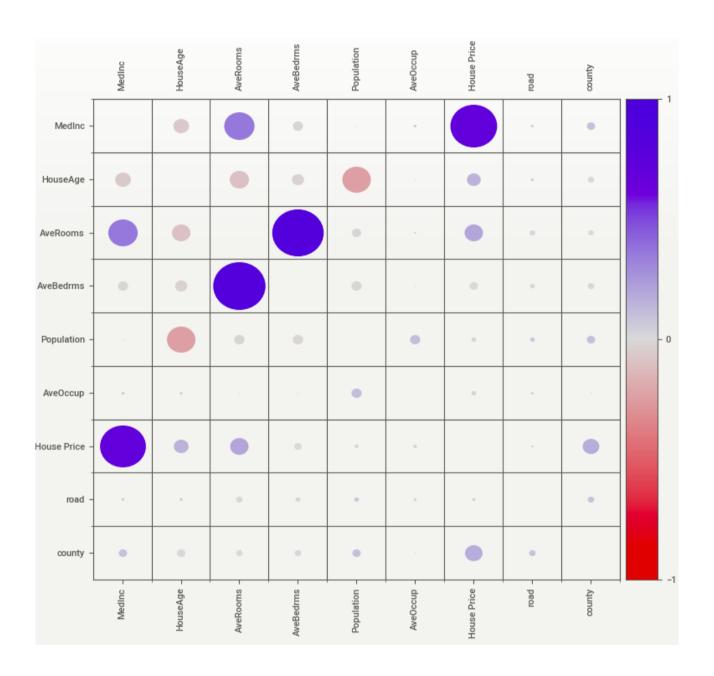
This is the cleaned/processed data which is now ready to be split into training and testing data and therafter to be fed into the ML algorithms





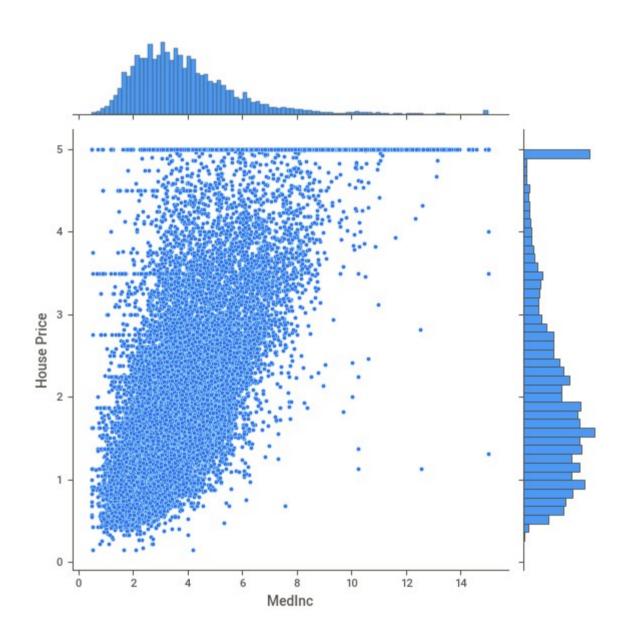
The above data gives us the max,min,null,distinct values present in each of the numerical attributes and more information like number of missing, Distinct, null values and their total sum etc.

Correlation Graph showing associations between various attributes



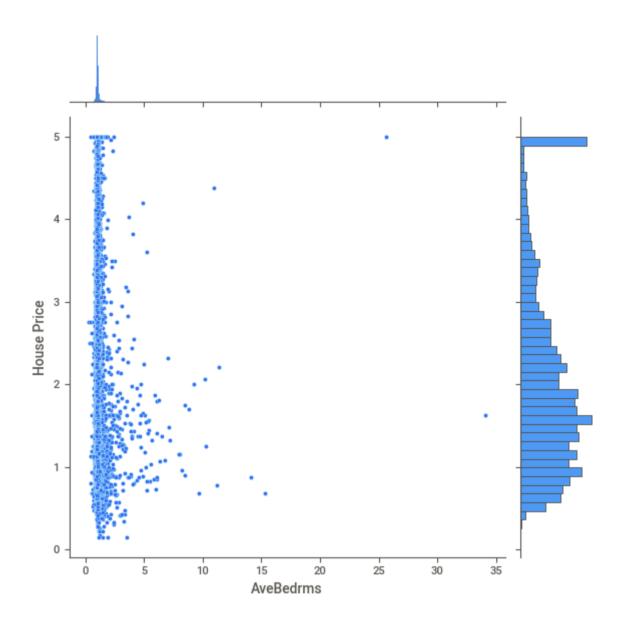
Circles are the symmetrical numerical correlations from -1 to 1.

Scatter plot showing the highest postive relationship out of all variables wrt to the target variable i.e "MedInc" and "House Price"



Scatter plot of target variable and "MedInc" attribute along with their respective histograms

Scatter plot showing the least postive relationship out of all variables wrt to the target variable i.e "AveBedrms" and "House Price"



Scatter plot of target variable and "AveBedrms" attribute along with their respective histograms

FINAL DATASET

```
df4.info()
<class 'pandas.core.frame.DataFrame'>
   Int64Index: 20640 entries, 18991 to 20555
   Data columns (total 8 columns):
        Column Non-Null Count
                                   Dtype
        MedInc
                    20640 non-null
    0
                                   float64
    1
        AveRooms
                   20640 non-null float64
    2
        AveBedrms
                   20640 non-null float64
        Population 20640 non-null float64
    3
    4
        Ave0ccup
                    20640 non-null float64
    5
        House Price 20640 non-null float64
                    20640 non-null int64
        road
        county 20640 non-null int64
    7
   dtypes: float64(6), int64(2)
   memory usage: 2.0 MB
```

With this the dataset is ready to be split into test and validate sets and therafter to be evaluated by the ML algorithm for model generation.

PYTHON PACKAGES IMPORTED

```
### Generic libraries
from geopy.geocoders import Nominatim
import matplotlib.pyplot as plt
import pandas as pd
import numpy as np
import seaborn as sns
import sweetviz as sv
import pickle
### Data set
from sklearn.datasets import fetch california housing
### sklearn modules
from sklearn.linear model import SGDClassifier
from sklearn.model selection import train_test_split
from sklearn.metrics import r2 score
from sklearn.metrics import mean squared error as MSE
from sklearn.model selection import KFold
from sklearn.metrics import r2 score
### ML algos
from sklearn.ensemble import RandomForestRegressor
from sklearn.neighbors import KNeighborsRegressor
from sklearn.linear model import LinearRegression
```

Sweetviz and **Seaborn** were used for data analytics and plotting the graphs along with **matploitlib**

geopy was used to get the address i.e the "road" and "county" attributes from the "Latitude" and "Longitude" values.

Sklearn was used for importing the necessary dataset as well as the ml aglorithms

PRE PROCESSING

Step 1 - Normalization

Here the features are rescaled to hard and fast range of [0,1] by subtracting minimum value of the feature then dividing by range.

```
x = df1.iloc[:,1:]

# Normalize values
df1.iloc[:,1:] = (x-x.min())/ (x.max() - x.min())
```

Step 2 - Splitting the dataset and standardization

```
from sklearn.preprocessing import StandardScaler

X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.2, random_state=101)

s_scaler = StandardScaler()

X_train = s_scaler.fit_transform(X_train.astype(np.float))

X_test = s_scaler.transform(X_test.astype(np.float))
```

The dataset is split into training and test datasets repectively in a 80:20 split with 80% training data and 20% test

Step 3 - Choosing the Model

Three models were chosen and all of them were evaluated for determining the best one

- K Neighbour Regression (*implemented from scratch*)
- Random Forest
- Linear Regression

ABOUT THE ALGOS

Random Forest:- utilizes ensemble learning, which is a technique that combines many classifiers to provide solutions to complex problems.

A random forest algorithm consists of many decision trees. The 'forest' generated by the random forest algorithm is trained through bagging.

Establishes the outcome based on the predictions of the decison trees. It predicts by tkaing the the aeverage or the mean of the output from various trees.

Linear Regression:- models the relationship between two variables by fitting a linear equation to the given data. One variable is considered to be an explanatory variable, and the other is considered to be a dependent variable

A linear regression line has an equation of the form Y = a + bX, where **X** is the explanatory variable and **Y** is the dependent variable.

The slope of the line is \boldsymbol{b} , and \boldsymbol{a} is the intercept (the value of \boldsymbol{y} when $\boldsymbol{x}=0$).

K Neighbour Regression :- This is a Regression algo based on k-nearest neighbors.

It has no model other than storing the entire dataset, so there is no learning required.

Predictions are made by searching through the entire training set for the K most similar neighbors and summarizing the output variable for those K instances. In case of regressin the it is usually the mean output variable

To measure the level of similarity a distance measure is used. For real-valued input variables, the most popular distance measure is Euclidean distance.

This algorithm has been implemented from scratch and the code for same is below

KNR IMPLEMENTATION

```
class KNeighbourRegressor():
  def init (self,k):
    self.k = k
  def fit(self,X,y):
    self.X train = X
    self.y train = y
  def predictions(self,x):
    distances = [euclid distance(x,x train) for x train in self.X train]
    k indices = np.argsort(distances)[:self.k]
   k nearest labels = [self.y train[i] for i in k indices]
    y pred = np.mean(k nearest labels)
    return y pred
  def predict(self,X):
    predicted labels = [self.predictions(x) for x in X]
    return np.array(predicted labels)
def euclid distance(a,b):
  dist = (a-b)**2
  return np.sqrt(np.sum(dist))
```

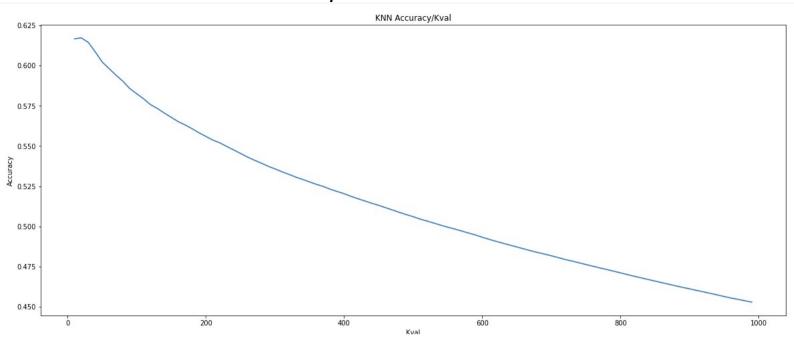
MODEL EVALUATION

Choosing the best K value

The following code and the graphs demonstrated how the correct value of K was achieved

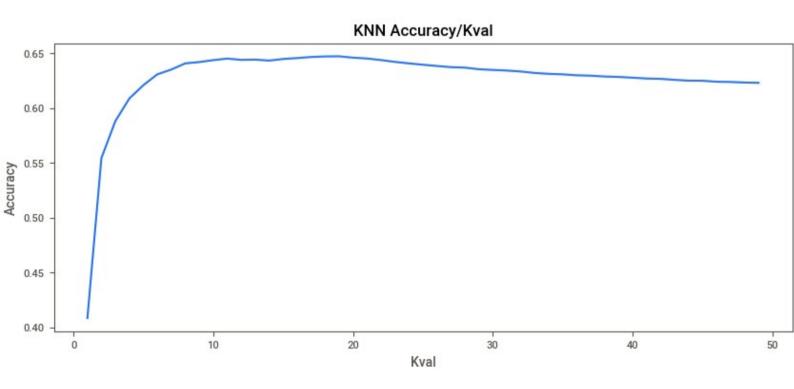
```
TestKvals = pd.DataFrame()
tmp = \{\}
for i in range(high,low):
  Kval = i
  tmp["Kval"] = Kval
  model = KNeighborsRegressor(Kval)
 model.fit(X_train, y_train)
  y pred = model.predict(X test)
 tmp["R2 val"] = (r2 score(y test, y pred))
  TestKvals = TestKvals.append([tmp])
TestKvals.set_index("Kval", inplace = True)
fig, axes = plt.subplots(ncols=1, figsize=(10,4))
TestKvals.R2 val.plot(ax=axes, kind="line", title="R2 Score")
plt.ylabel('Accuracy')
plt.title('KNN Accuracy/Kval')
plt.show()
```

K from 0 to 1000



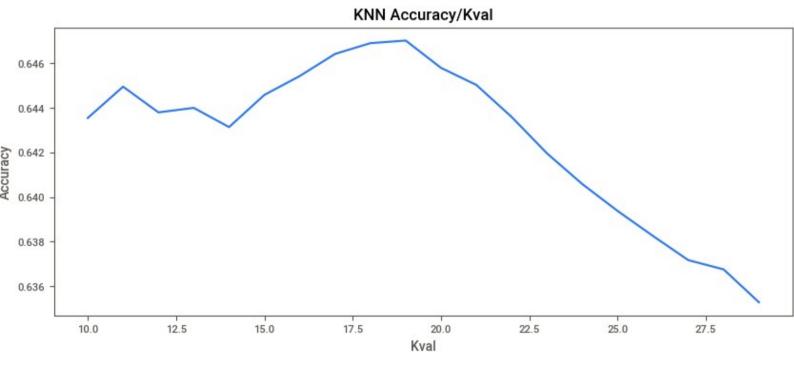
As we see here the best K value somewhere between 0 and 50 approximately

K from 0 to 50



We see the required K value is in the range of 10 to 30 approximately with value being very close to 20

K from 10 to 30



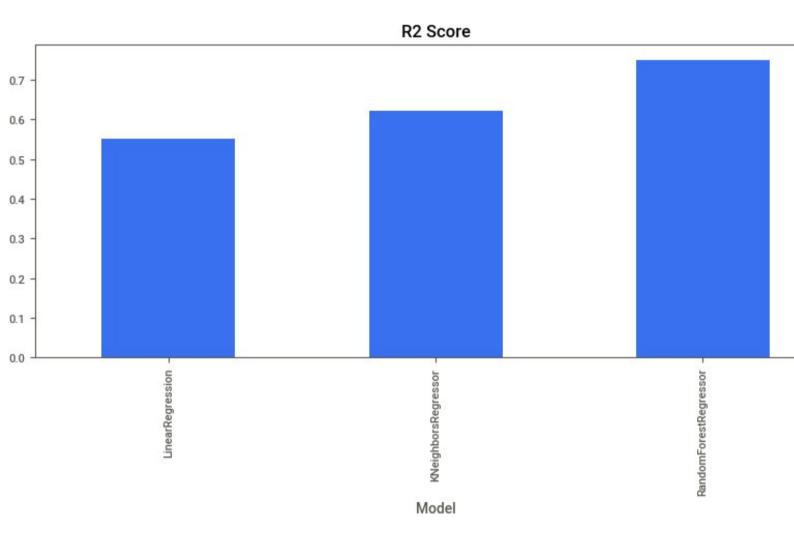
Hence the required K value is between 17.5 and 20 and can be approximated to 19 or 20

PLOTTING ACCURACIES OF ALL THREE ALGORITHMS

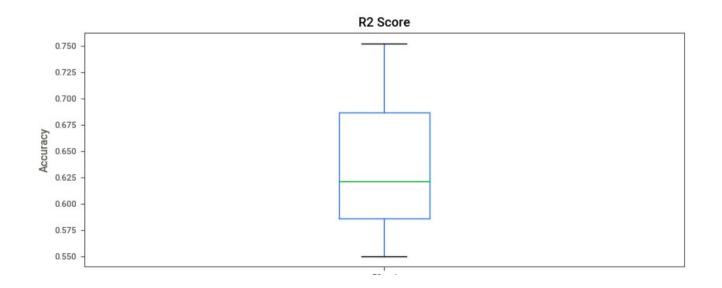
Here all the three algorithms are evaluated and their repective accuracies are observed and plotted to determine the best algorithm out of the three

```
models = [ LinearRegression(), KNeighborRegressor(20), RandomForestRegressor() ]
TestModels = pd.DataFrame()
tmp = \{\}
for model in models:
  print(model)
  m name = str(model)
  tmp["Model"] = m name[:m name.index("(")]
  model.fit(X train, y train)
  y pred = model.predict(X test)
  tmp["R2 val"] = (r2 score(y test, y pred))
  TestModels = TestModels.append([tmp])
TestModels.set index("Model", inplace = True)
fig, axes = plt.subplots(ncols=1, figsize=(10,4))
TestModels.R2 val.plot(ax=axes, kind="line", title="R2 Score")
plt.ylabel('Accuracy')
plt.show()
```

Bar Graph depicting the most accurate algorithm for the model



Box plot of accuracies and algorithms



Since highest accuracy is of Random Forest we predict the values using Random forest itself

PREDICTING THE VALUES

```
# model prediction

y_pred = model.predict(X_test)
print(y_pred)

[0.27038875 0.54496724 0.46170559 ... 0.41702539 0.46172416 0.59865571]
```

Creating Dataframe

```
dframe = pd.DataFrame({'Actual':y_test.flatten(),'Predicted':y_pred.flatten()})
dframe.head(25)
```

MODEL ACCURACY

Accuracy of the model

```
MSE_score = MSE(y_test,y_pred)

print('R2 score', r2_score(y_test, y_pred))
print('Mean Absolute Error:', mean_absolute_error(y_test, y_pred))
print('Mean Squared Error:', MSE_score.mean())
print('Root Mean Squared Error:', np.sqrt(MSE_score))
```

```
R2 score 0.7661347030758191
Mean Absolute Error: 0.08079751176869304
Mean Squared Error: 0.01399506695335315
Root Mean Squared Error: 0.11830074789853676
```

Visualizing our model's predictions

